

# Creating Deployments

---



**Dan Wahlin**

WAHLIN CONSULTING

@danwahlin [www.codewithdan.com](http://www.codewithdan.com)



# Module Overview

Deployments Core Concepts

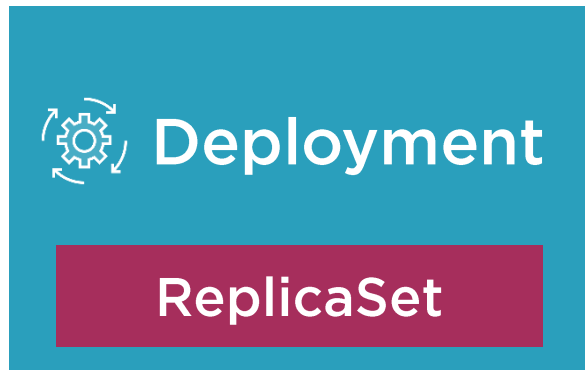
Creating a Deployment

kubectl and Deployments



Deployment Options



# You Are Here



# You Are Here



# Deployments Core Concepts

---



A ReplicaSet is a declarative way to manage Pods



A Deployment is a declarative way to manage Pods using a ReplicaSet




**Pods represent the most basic resource in Kubernetes**

**Can be created and destroyed but are never re-created**

**What happens if a Pod is destroyed?**


**Deployments and ReplicaSets ensure Pods stay running and can be used to scale Pods**

Pods,  
Deployments,  
and ReplicaSets





## The Role of ReplicaSets



### ReplicaSets act as a Pod controller:

- Self-healing mechanism
- Ensure the requested number of Pods are available
- Provide fault-tolerance
- Can be used to scale Pods
- Relies on a Pod template
- No need to create Pods directly!
- Used by Deployments

# Result of Creating a ReplicaSet

```
iMac-3:replicaSets danwahlin$ k get all
```

```
+ kubectl get all
```

| NAME               | READY | STATUS  | RESTARTS | AGE |
|--------------------|-------|---------|----------|-----|
| pod/frontend-8rslg | 1/1   | Running | 0        | 8s  |
| pod/frontend-zc8gl | 1/1   | Running | 0        | 8s  |

| NAME               | TYPE      | CLUSTER-IP | EXTERNAL-IP | PORT(S) | AGE |
|--------------------|-----------|------------|-------------|---------|-----|
| service/kubernetes | ClusterIP | 10.96.0.1  | <none>      | 443/TCP | 41d |

| NAME                     | DESIRED | CURRENT | READY | AGE |
|--------------------------|---------|---------|-------|-----|
| replicaset.apps/frontend | 2       | 2       | 2     | 8s  |


2 Pods created



ReplicaSet created



# The Role of Deployments



## A Deployment manages Pods:

- Pods are managed using ReplicaSets
- Scales ReplicaSets, which scale Pods
- Supports zero-downtime updates by creating and destroying ReplicaSets
- Provides rollback functionality
- Creates a unique label that is assigned to the ReplicaSet and generated Pods
- YAML is very similar to a ReplicaSet

# Creating a Deployment

---




# Defining a Deployment with YAML



Deployment

+ kubectl =



## Defining a Deployment (From a High-Level)

```
apiVersion: apps/v1
kind: Deployment
metadata:
spec:
  selector:
  template:
    spec:
      containers:
      - name: my-nginx
        image: nginx:alpine
```

- ◀ Kubernetes API version and resource type (Deployment)
- ◀ Metadata about the Deployment
- ◀ Select Pod template label(s)
- ◀ Template used to create the Pods
- ◀ Containers that will run in the Pod



# Defining a Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend
  labels:
    app: my-nginx
    tier: frontend
spec:
  selector:
    matchLabels:
      tier: frontend
  template:
    metadata:
      labels:
        tier: frontend
    spec:
      containers:
        - name: my-nginx
          image: nginx:alpine
```

- ◀ Kubernetes API version and resource type (Deployment)
- ◀ Metadata about the Deployment
- ◀ The selector is used to "select" the template to use (based on labels)
- ◀ Template to use to create the Pod/Containers (note that the selector matches the label)



```
apiVersion: apps/v1

kind: Deployment

...

template:
  spec:
    containers:
    - name: my-nginx
      image: nginx:alpine
      livenessProbe:
        httpGet:
          path: /index.html
          port: 80
        initialDelaySeconds: 15
        timeoutSeconds: 2
        periodSeconds: 5
        failureThreshold: 1
```

- ◀ Define liveness probe (readiness probes can also be defined)
- ◀ Check /index.html on port 80





# kubectl and Deployments

---



## # Create a Deployment

```
kubectl create -f file.deployment.yml
```

### Creating a Deployment

Use the **kubectl create** command along with the **--filename** or **-f** switch



# Creating or Applying Changes

Use the **kubectl apply** command along with the **--filename** or **-f** switch

```
# Alternate way to create or apply changes to a  
# Deployment from YAML
```

```
kubectl apply -f file.deployment.yml
```

```
# Use --save-config when you want to use  
# kubectl apply in the future
```

```
kubectl create -f file.deployment.yml --save-config
```

Store current  
properties in  
resource's annotations



```
kubectl get deployments
```

Getting Deployments

**List all Deployments**



# List all Deployments and their labels

```
kubectl get deployment --show-labels
```

# Get all Deployments with a specific label

```
kubectl get deployment -l app=nginx
```

## Deployments and Labels

List the labels for all Deployments using the **--show-labels** switch

To get information about a Deployment with a specific label, use the **-l** switch



# Deleting a Deployment

To delete a Deployment use **kubectl delete**

Will delete the Deployment and all associated Pods/Containers

```
# Delete Deployment
```

```
kubectl delete deployment [deployment-name]
```

# Scale the Deployment Pods to 5

```
kubectl scale deployment [deployment-name] --replicas=5
```

# Scale by referencing the YAML file

```
kubectl scale -f file.deployment.yml --replicas=5
```

## Scaling Pods Horizontally

Update the YAML file or use the **kubectl scale** command

```
spec:  
  replicas: 3  
  selector:  
    tier: frontend
```



# kubectl Deployments in Action

---





```
kubectl create -f nginx.deployment.yml --save-config  
kubectl describe [pod | deployment] [pod-name | deployment-name]  
kubectl apply -f nginx.pod.yml  
kubectl get deployments --show-labels  
kubectl get deployments -l app=my-nginx  
kubectl scale -f nginx.deployment.yml --replicas=4
```

## kubectl Deployments Commands

**Several different kubectl commands can be used to create and work with Deployments**




# Deployment Options

---




# How Do You Update Existing Pods?



Zero downtime deployments  
allow software updates to be  
deployed to production without  
impacting end users



## Deployment Options



One of the strengths of Kubernetes is zero downtime deployments


Update an application's Pods without impacting end users

Several options are available:

- Rolling updates
- Blue-green deployments
- Canary deployments
- Rollbacks


# Rolling Deployments

## Initial Pod State




# Rolling Deployments

## Rollout New Pod



# Rolling Deployments


## Delete Pod






# Rolling Deployments

## Rollout New Pod




# Rolling Deployments

## Delete Pod






# Rolling Deployments

## Rollout New Pod




# Rolling Deployments

## Delete Pod



# Rolling Deployments

## Rollout New Pod



# Updating a Deployment

Update a deployment by changing the YAML and applying changes to the cluster with **kubectl apply**

```
# Apply changes made in a YAML file  
kubectl apply -f file.deployment.yml
```

# Zero Downtime Deployments in Action

---



# Summary



Pods are deployed, managed, and scaled using deployments and ReplicaSets

Deployments are a higher-level resource that define one or more Pod templates

The `kubectl create` or `kubectl apply` commands can be used to run a deployment

Kubernetes supports zero downtime deployments

